

Line breaking

Victor Eijkhout

August 2004

Breaking a paragraph into lines is the problem of, given a string of words and other material with intervening spaces, breaking that string into chunks (lines) of approximately equal length, and doing so in a visually attractive way. Simple strategies (see the ‘first fit’ algorithm below) give a result that is easy to compute, but that can be visually very unappealing. While the problem of finding globally optimal line breaks sounds very hard – with n words there are 2^n ways of breaking the paragraph; also, this problem resembles the bin-packing problem which is NP-complete – it can actually be solved fairly efficiently.

T_EX’s basic strategy is to calculate the badness of breaking the lines at certain points, and to minimize the badness over the whole paragraph.

1 The elements of a paragraph

T_EX’s paragraph breaking algorithm is based around the concepts of

- Boxes: this comprises letters, formulas, T_EX boxes, and other material of a fixed width.
- Glue: this is white space; a glue item has a natural width, stretchability, and shrinkability.
- Penalties: these are items that express the desirability or undesirability of breaking a line at a particular point.

The same elements are also present in a vertical list; in both cases there are some other, more rare items, that we will ignore here.

1.1 Some details

1.1.1 Boxes

The boxes in a paragraph are such things as words, rules, math formulas, and actual T_EX `\boxes`. A box can not be broken: it is completely described by its height, depth, width. Actually, looking at words as boxes is too simplistic, since words can often be hyphenated. This means that a word is a sequence of boxes alternating with penalties.

1.1.2 Penalties

A penalty item describes the desirability or undesirability of breaking a list at some point. Penalties can be inserted by hand (most often as the `\nobreak` macro, which is equivalent to `\penalty10000`), or in a macro, or are inserted by \TeX itself. An example of the latter case is the `\hyphenpenalty` which discourages breaking at a hyphen in a word.

Hyphenating a word can be simulated by having a penalty of zero at the hyphenation location. Since this usually introduces a hyphen character, \TeX internally pretends that penalties can have a width if the list is broken at that place.

The most common types of penalty are the infinite penalty that starts a non-breaking space, and the penalty associated with breaking by hyphenating a word. The latter kind is called a ‘flagged penalty’, and \TeX has an extra amount of demerits associated with them, for instance to prevent two consecutive lines ending in a hyphen.

Penalties can have positive and negative values, to discourage or encourage breaking at a specific place respectively. The values $+\infty$ and $-\infty$ are also allowed, corresponding to a forbidden and forced break respectively.

1.1.3 Glue

A ‘glue’ is a dimension with possible stretch and/or shrink. There are glue denotations, such as `2cm plus .5cm minus .1cm`, or glue parameters, such as `\leftskip` or `\abovedisplayskip`. The parameters are inserted automatically by the various \TeX mechanisms.

Glue can be measured in points `pt`, centimeters `cm`, millimeters `mm`, inches `in`. There is also infinite glue: `fil`, `fill`, and `filll`. Presence of \TeX ’s infinite glue (`fill`) causes all other glue to be set at natural width, that is, with zero stretch and shrink.

If there is more than one glue item in a list, the natural widths and the stretch and shrink amounts are added together. This means that a list with both `2cm plus 1cm` and `2cm plus -1cm` has no stretch since the stretch amounts add up to zero. On the other hand, with `2cm plus 1cm` and `2cm minus 1cm` it has both stretch and shrink.

The stretch and shrink components of glue are not treated symmetrically. While in a pinch we can allow a line to be stretched further than the indicated maximum, we can not allow spaces to be shrunk to zero, or even close to that.

1.1.4 Stretch and shrink

Each space can have stretch and shrink. When we consider a line, we add up all the stretch and shrink and compute an ‘adjustment ratio’ as the ratio of the shortfall or excess space to the available stretch or shrink respectively. This ratio r is negative for lines that need to be shrunk.

A simple algorithm would be to impose a limit of $|r| \leq 1$ (and then to minimize the number of hyphenations under that restriction), but that might be too restrictive. Instead, \TeX uses a concept of ‘badness’. The badness of a line break is infinite if $r < -1$; otherwise it is cubic in the absolute size of r .

1.1.5 Line break locations

Here are the main (but not the only) locations where \TeX can decide to break a line.

- At a penalty
- At a glue, if it is preceded by a non-discardable item, meaning, not a penalty or other glue
- At a hyphen in a word
- At a place where \TeX knows how to hyphenate the word. (There is actually a mechanism, called ‘discretionaries’ that handles these last two cases.)

1.2 Examples

Here are a few examples of the things that the boxes/penalties/glue mechanism is capable of.

1.3 Centered text

By using `\leftskip` and `\rightskip` we can get centered text.

```
\begin{minipage}{4cm}
\leftskip=0pt plus 1fil \rightskip=0pt plus 1fil
\parfillskip=0pt
This paragraph puts infinitely stretchable glue at
the left and right of each line.
The effect is that the lines will be centered.
\end{minipage}
```

Output:

This paragraph puts
infinitely stretchable glue at
the left and right of each
line. The effect is that the
lines will be centered.

The following centers only the last line. This is done by letting the `\leftskip` and `\rightskip` cancel each other out, except on the last line.

```
\begin{minipage}{5cm}
\leftskip=0pt plus 1fil \rightskip=0pt plus -1fil
\parfillskip=0pt plus 2fil
This paragraph puts infinitely stretchable glue at
the left and right of each line, but the amounts cancel out.
The parfillskip on the last line changes that.
\end{minipage}
```

Output:

This paragraph puts infinitely
stretchable glue at the left and right
of each line, but the amounts cancel
out. The parfillskip on the last line
changes that.

1.3.1 Hanging punctuation

Hanging punctuation is a style of typesetting where punctuation that would wind up against the right margin is actually set *in* the right margin. Optically, this makes the margin look straighter.

```
\newbox\pbox \newbox\cbox
\setbox\pbox\hbox{.} \wd\pbox=0pt
\setbox\cbox\hbox{,} \wd\cbox=0pt
\newdimen\csize \csize=\wd\cbox
\newdimen\psize \psize=\wd\pbox

\catcode`,=13 \catcode`=13
\def,{\copy\cbox \penalty0 \hskip\csize\relax}
\def. {\copy\pbox \penalty0 \hskip\psize\relax}
```

Here is a bit of text
in a minipage, with
too much punctuation.
Every sentence, long
or short, is overly, yes,
overly, punctuated. This
should show off ‘hanging
punctuation’.

1.3.2 Mathematical Reviews

In ‘Mathematical Reviews’ the name of the reviewer should be separated sufficiently from the review, but fit on the same line if space allows.

This review is rather negative, almost
devastating. A. Reviewer

This review is als very negative but it’s
longer than the other.
A. Nother-Reviewer

We do this by having two separate infinite glues with a break in between, and with a total natural width equal to the minimum separation. The trick is to make sure that the second glue is not discarded after a break, which we do by putting an invisible box at the beginning.

```
\def\signed#1{\unskip
\penalty10000 \hskip 40pt plus 1fill
\penalty0
\hbox{}\penalty10000
\hskip 0pt plus 1fill
\hbox{#1}%
\par
}
```

2 T_EX's line breaking algorithm

2.1 Elements

2.1.1 Glue setting and badness

In order to make a list fit a prescribed dimension, there is a process called ‘glue setting’. The natural size of the list and the desired size are compared. Let ρ be the ratio of the amount stretched (shrunk) to the possible amount of stretch (shrink). The exact definition is such that the ratio is positive for stretch and negative for shrink: let ℓ be the desired length of the line, L the natural width, X the available stretch and Y the available shrink, then

$$\rho = \begin{cases} 0 & \ell = L \\ (\ell - L)/X & \text{(stretch:) } \ell > L \text{ and } X > 0 \\ (\ell - L)/Y & \text{(shrink:) } \ell < L \text{ and } Y > 0 \\ \text{undefined} & \text{otherwise} \end{cases}$$

Then the ‘badness’ of the needed glue setting is

$$b = \begin{cases} 10\,000 & \rho < 1 \text{ or undefined} \\ \min \{10\,000, 100|\rho|^3\} & \text{otherwise} \end{cases}$$

Since 10 000 is considered infinite in glue arithmetic, this algorithm allows glue to be stretched further than the indicated amount, but not shrunk beyond what is available.

A list that is stretched or shrunk is put in one of the following four categories:

tight (3) if it has shrunk with $b \geq 13$

decent (2) if $b \leq 12$

loose (1) if it has stretched with $100 > b \geq 13$

very loose (0) if it has stretched with $b \geq 100$

Note that $100 \times (1/2)^3 = 12.5$, so the crossover values denote that half the stretch or shrink is used.

Lines that differ by more than one in their classifications are called ‘visually incompatible’.

2.1.2 Demerits

Breaking a line at a certain points gives the penalty p associated with that point, and the badness b of the resulting stretch/shrink. These are combined into a ‘demerits’ figure:

$$d = \begin{cases} b^2 + p^2 & 0 \leq p < 10\,000 \\ b^2 - p^2 & -10\,000 < p < 0 \end{cases}$$

The demerits for breaking a paragraph along a certain sequence of break points is then the sum of the demerits of the lines, plus `\adjdemerits` for every two lines that are not visually compatible (section ??), `\doublehyphendemerits` for pairs of lines that end in a hyphen, and `\finalhyphendemerits` if the last full line ends in a hyphen.

T_EX acts as if before the first line there is a line that is ‘decent’; the last line will typically contain infinite glue, so all spaces are set at natural width.

For full generality, the last line of a paragraph is handled like any other. Filling out the line to the margin is realized by added infinite glue plus a trailing penalty that forces a line break at the end of the paragraph.

2.2 Breaking strategies

We will now look at a couple of line breaking strategies. The first two will be strictly local; the third – \TeX 's algorithm – is able to optimize in a more global sense.

The problem with local algorithms is that they can not take a slightly worse solution in one line to prevent much worse from happening later. This will for instance allow tight and very loose lines to occur next to each other.

2.2.1 First fit

The traditional algorithm for line breaking only looks at the current line. Once a word is starting to exceed the right margin, the following cases are investigated.

1. If the spaces in the line can be compressed without exceeding some maximum shrinkage, break after this word.
2. Otherwise, if the spaces can be stretched short of some maximum, break before this word.
3. Otherwise, try hyphenating this word.
4. If no hyphenation point can be found, accept a line with spaces stretched to whatever degree is needed to break before this word.

If you have set text with \TeX , you will have noticed that \TeX 's solution to the last point is slightly different. It will let a word protrude into the right margin as a visual indicator that no good breakpoint could be found. (\TeX 's tolerance to bad breaks can be increased by setting the `\emergencystretch` parameter.)

This method can be called 'first fit', because it will the first option (compress), without comparing if later options (stretching) look better. This is remedied in \TeX by, instead of having an all-or-nothing it fits / it doesn't fit distinction, there is a continuous scale of evaluation.

2.2.2 Best fit

A slight improvement on the first fit algorithm results from deciding between the possibilities 1–3 based on badness calculations. This can be termed 'best fit', and while it may work slightly better than fit, it is still a local decision process.

2.2.3 Total fit

\TeX 's actual algorithm calculates the 'demerits' of a line break as a compound of badness, the breakpoint penalty, plus any flagged penalties. It then adds together the demerits of the whole paragraph, and minimizes this number. This makes it possible to use a slightly worse line break early in the paragraph, to prevent a much worse one later.

Exercise 1. In dynamic programming, many solutions start from a final stage and work backwards. Why is this approach inappropriate for \TeX 's line breaking algorithm? Why would it be even less appropriate for a page breaking algorithm?

2.3 Model implementation

We will here only discuss implementations of solutions based on dynamic programming.

The line breaking algorithm goes linearly through the items in the horizontal list, and for each considers whether there is a valid breakpoint after it, and with what cost. For the latter point, it needs to know what the beginning of the line is, so there is an inner loop over all preceding words. This would make the running time of the algorithm quadratic in the number of words, which is much better than the initial estimate of 2^n .

However, the situation is better than that. The number of words that can fit on a line is limited by what can fit when all spaces are squeezed to zero. The easiest way to implement this is by maintaining an ‘active list’ of possible words to begin the line with. A word can be removed from the active list if the material from it to the current word does not fit: if it does not fit now, it will certainly not fit when the next word is considered.

This is then the main program; we will mainly vary the function that computes the breakpoint cost.

```
active = [0]
nwords = len(paragraph)
for w in range(1,nwords):
    # compute the cost of breaking after word w
    for a in active:
        line = paragraph[a:w+1]
        ratio = compute_ratio(line)
        if w==nwords-1 and ratio>0:
            ratio = 0 # last line will be set perfect
        print "..line=",line,"; ratio=",ratio
        if ratio<-1:
            active.remove(a)
            print "active point",a,"removed"
        else:
            update_cost(a,w,ratio)
    report_cost(w)
    active.append(w)
    print
```

The only thing different between various strategies is how the cost of a line break is computed by `update_cost(a,w,ratio)`.

Exercise 2. Not every word should be added to the active list. For instance, for any realistic line length, the second word in the paragraph will not have a valid breakpoint after it, so we need not consider it. Take the model implementation and add this modification. Measure the improvement in running time, for instance by counting the number of calls to some inner routine. Give a theoretical argument for how this reduces the complexity of the algorithm.

2.3.1 First fit implementation

Since at first we are looking only locally, for each breakpoint we only keep track of the cost and the previous breakpoint that the cost was computed from. Here we set up the data structure `cost`. Element `cost[w]` describes the cost of breaking after word `w`; the 'from' component is the number of the first word of the line.

```
def init_costs():
    global cost
    cost = len(paragraph)*[0]
    for i in range(len(paragraph)):
        cost[i] = {'cost':0, 'from':0}
    cost[0] = {'cost':10000, 'from':-1}
```

The essential function is the cost computation. In first fit we accept any stretch or shrink that is $|\rho| < 1$.

```
def update_cost(a,w,ratio):
    global cost
    if a>0 and cost[a-1]['cost']<10000:
        if ratio<=1 and ratio>=-1:
            to_here = abs(ratio)
        else: to_here = 10000
        if cost[w]['cost']==0 or to_here<cost[w]['cost']:
            cost[w]['cost'] = to_here; cost[w]['from'] = a-1
```

(The first test serves to make sure that the previous point being considered is in fact a valid breakpoint.)

Here is the remaining function that constructs the chain of breakpoints:

```
def final_report():
    global cost,nwords,paragraph
    print "Can break this paragraph at cost",\
        cost[nwords-1]['cost']
    cur = len(paragraph)-1; broken = []
    while cur!=-1:
        prev = cost[cur]['from']
        line = paragraph[prev+1:cur+1]
        broken.insert(0,line)
        cur = prev;
    set_paragraph(broken)
```

A small example text, faked in monospace:

You may never have thought of it, but fonts (better:	-0.111111111111
typefaces) usually have a mathematical definition somehow.	-0.666666666667
If a font is given as bitmap, this is often	0.888888888889
a result originating from a more compact description.	0.0
Imagine the situation that you have bitmaps at 300dpi, and	-0.777777777778
you buy a 600dpi printer. It wouldn't look pretty.	0.25
There is then a need for a mathematical way of	0.555555555556
describing arbitrary shapes. These shapes can also be	0.0
three-dimensional; in fact, a lot of the mathematics in	-0.285714285714


```
this chapter was developed by a car manufacturer for 0.0
modeling car body shapes. But let us for now only 0.222222222222
look in two dimensions, which means that the curves 0.125
are lines, rather than planes.
```

We see ugly stretched break in line 3, especially after the compressed line 2. However, both of them fit the test.

It is in fact simple to turn this into a dynamic programming solution that considers a global minimum:

```
def update_cost(a,w,ratio):
    global cost
    if ratio<=1 and ratio>=-1:
        to_here = abs(ratio)
    else: to_here = 10000
    if a>0:
        from_there = cost[a-1]['cost']
        to_here = to_here+from_there
    else: from_there = 0
    if cost[w]['cost']==0 or to_here<cost[w]['cost']:
        cost[w]['cost'] = to_here; cost[w]['from'] = a-1
```

2.3.2 Best fit

In the best fit strategy, we compute badness from the stretch/shrink ratio. This involves only a slight change in the cost computation function:

```
def update_cost(a,w,ratio):
    global cost
    to_here = 100*abs(ratio)**2
    if a>0:
        from_there = cost[a-1]['cost']
        to_here = to_here+from_there
    else: from_there = 0
    if cost[w]['cost']==0 or to_here<cost[w]['cost']:
        cost[w]['cost'] = to_here; cost[w]['from'] = a-1
```

The same text:

```
You may never have thought of it, but fonts (better: -0.111111111111
typefaces) usually have a mathematical definition somehow. -0.666666666667
If a font is given as bitmap, this is often a 0.5
result originating from a more compact description. 0.5
Imagine the situation that you have bitmaps at 300dpi, and -0.777777777778
you buy a 600dpi printer. It wouldn't look pretty. 0.25
There is then a need for a mathematical way of 0.555555555556
describing arbitrary shapes. These shapes can also be 0.0
three-dimensional; in fact, a lot of the mathematics in -0.285714285714
this chapter was developed by a car manufacturer for 0.0
modeling car body shapes. But let us for now only 0.222222222222
look in two dimensions, which means that the curves 0.125
are lines, rather than planes.
```

While there are no lines stretched with $\rho >$, the quadratic function has improved the break in line 3.

2.3.3 Total fit

For the algorithm that \TeX uses, we have to distinguish between lines that are tight, decent, loose. This makes our datastructure more complicated:

```
def init_costs():
    global cost
    nul = [0,0,0]
    cost = len(paragraph)*[ 0 ]
    for i in range(len(paragraph)):
        cost[i] = nul[:]
        for j in range(3):
            cost[i][j] = {'cost':10000, 'from':-2}
    for j in range(3):
        cost[0][j] = {'cost':10000, 'from':-1}
```

An element `cost[i]` is now an array of three possible breakpoints, one in each of the classifications. An actual breakpoint is now in `cost[word][type]['from']` and `cost[word][type]['cost']`.

The cost computation becomes more complicated:

```
def minimum_cost_and_type(w):
    global cost
    c = 10000; t = 0
    for type in range(3):
        nc = cost[w][type]['cost']
        if nc < c:
            c = nc; t = type
    return [c,t]
def update_cost(a,w,ratio):
    global cost
    type = stretch_type(ratio)
    to_here = 100*abs(ratio)**2
    if a > 0:
        [from_there,from_type] = minimum_cost_and_type(a-1)
        to_here += from_there
    else: from_there = 0
    if cost[w][type]['cost'] == 0 or \
        to_here < cost[w][type]['cost']:
        cost[w][type]['cost'] = to_here;
        cost[w][type]['from'] = a-1
```

Exercise 3. The total fit code does not yet contain the equivalent of \TeX 's `\adjdemerits`. Add that.

Let us look at the same test text again:

You may never have thought of it, but fonts (better: typefaces) usually have a mathematical definition somehow. If a font is given as bitmap, this is often a result originating from a more compact description. Imagine the situation that you have bitmaps at 300dpi, and you buy a 600dpi printer. It wouldn't look pretty. There is then a need for a mathematical way of describing arbitrary shapes. These shapes can also be three-dimensional; in fact, a lot of the mathematics in this chapter was developed by a car manufacturer for modeling car body shapes. But let us for now only look in two dimensions, which means that the curves are lines, rather than planes.

```
-0.111111111111
1.2
-0.454545454545
0.5
1.0
-0.333333333333
-0.4
0.0
-0.285714285714
0.0
0.222222222222
0.125
```

In this output, line 2 is stretched further than before, to prevent lower badnesses later.

Exercise 4. Add the functionality for hanging indentation to this code.

Exercise 5. (bonus point exercise) \TeX has the possibility of forcing a paragraph to be a line longer or shorter than optimal. Implement that.

2.3.4 Utility parts

File header: we read a text and store it.

```
#!/usr/bin/env python

import sys

max_line_length = 60

paragraph = []
while 1:
    try:
        a = raw_input()
        paragraph.extend(a.split())
    except (EOFError):
        break
```

In order to fake stretch and shrink with a monospace font, we let a 'space' be two spaces by default.

```
def line_length(words):
    l = 2*(len(words)-1)
    for w in words:
        l += len(w)
    return l

#
# ratio = -1 : shrink each double space to one
# ratio = 1 : stretch each double space to three
#
def compute_ratio(line):
    spaces = len(line)-1
```

```

    need = 1.*(max_line_length-line_length(line))
    #print "ratio:",need,spaces
    if spaces==0: return 10000
    else: return need/spaces

```

Output formatting with this idea:

```

def set_paragraph(para):
    for l in range(len(para)-1):
        line = para[l]
        set_line(line)
    set_last_line(para[len(para)-1])
def set_line(line):
    shortfall = max_line_length-line_length(line)
    for w in range(len(line)-1):
        sys.stdout.write(line[w])
        if shortfall>0:
            sys.stdout.write(' '); shortfall = shortfall-1
        elif shortfall<0:
            sys.stdout.write(' '); shortfall = shortfall+1
        else:
            sys.stdout.write(' ')
    sys.stdout.write(line[len(line)-1])
    print " ",compute_ratio(line)
def set_last_line(line):
    for w in range(len(line)-1):
        sys.stdout.write(line[w])
        sys.stdout.write(' ')
    sys.stdout.write(line[len(line)-1])
    print

```

Contents

1	Introduction	1	3.1	<i>Class P</i>	4
2	Basics	3	3.2	<i>Class NP</i>	4
2.1	<i>Optimization versus decision problems</i>	3	3.3	<i>Examples</i>	5
2.2	<i>Language of a problem</i>	3	4	NP-completeness	5
2.3	<i>Turing machines</i>	3	4.1	<i>Transformations</i>	5
3	Complexity classes	4	4.2	<i>NP-complete</i>	6
			4.3	<i>Proving NP-completeness</i>	6